

# Multi-GPU accelerated multi-spin Monte Carlo simulations of the 2D Ising model



# What was done?

- ▶ **Ising model**

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- ▶ **Metropolis algorithm:**

Efficient sampling

$$p_a = e^{-\frac{E}{k_B T}}$$

# What was done?

- ▶ **Ising model**

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- ▶ **Metropolis algorithm:**

Efficient sampling

$$p_a = e^{-\frac{E}{k_B T}}$$

Single core CPU



# What was done?

- **Ising model**

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- **Metropolis algorithm:**

Efficient sampling

$$p_a = e^{-\frac{E}{k_B T}}$$

Single core CPU



Single GPU



# What was done?

- **Ising model**

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- **Metropolis algorithm:**

Efficient sampling

$$p_a = e^{-\frac{E}{k_B T}}$$

Single core CPU



Single GPU



Multiple GPUs



# What was done?

- **Ising model**

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- **Metropolis algorithm:**

Efficient sampling

$$p_a = e^{-\frac{E}{k_B T}}$$

Single core CPU



Single GPU



Multiple GPUs



How does it scale? Is it worth the effort?

# Ising model I

- Formula:

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

# Ising model I

- ▶ Formula:

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- ▶ A standard model of statistical physics



# Ising model I

- ▶ Formula:

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$



- ▶ A standard model of statistical physics
- ▶ classical model

# Ising model I

- Formula:

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- A standard model of statistical physics
- classical model
- Describes magnets:





# Ising model I

- ▶ Formula:

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

- ▶ A standard model of statistical physics
- ▶ classical model
- ▶ Describes magnets:



- ▶ System sizes:  $100'000 \times 100'000$

## Ising model II


Nearest neighbour interactions only!

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

## Ising model II

Nearest neighbour interactions only!


$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$



## Ising model II

Nearest neighbour interactions only!

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$



Calculation of energy:  $\mathcal{O}(n^2) \Rightarrow \mathcal{O}(n)$

# Metropolis algorithm

- **Goal:** Sample phase space

$$\begin{array}{ccccc} \uparrow & \downarrow & \downarrow & \uparrow & \downarrow \\ \downarrow & \uparrow & \downarrow & \downarrow & \downarrow \\ \uparrow & \uparrow & \downarrow & \uparrow & \downarrow \\ \uparrow & \downarrow & \uparrow & \uparrow & \downarrow \end{array} \sim e^{-\frac{E}{k_B T}}$$

# Metropolis algorithm

- **Goal:** Sample phase space

$$\begin{array}{ccccc} \uparrow & \downarrow & \downarrow & \uparrow & \downarrow \\ \downarrow & \uparrow & \downarrow & \downarrow & \downarrow \\ \uparrow & \uparrow & \downarrow & \uparrow & \downarrow \\ \uparrow & \downarrow & \uparrow & \uparrow & \downarrow \end{array} \sim e^{-\frac{E}{k_B T}}$$

- **Algorithm:**



# Metropolis algorithm

- **Goal:** Sample phase space


↑	↓	↓	↑	↓
↓	↑	↓	↓	↓
↑	↑	↓	↑	↓
↑	↓	↑	↑	↓

$\sim$

$$e^{-\frac{E}{k_B T}}$$

- **Algorithm:**

1.) Propose new state: Random spin flips!



# Metropolis algorithm

- **Goal:** Sample phase space

$$\begin{array}{ccccc} \uparrow & \downarrow & \downarrow & \uparrow & \downarrow \\ \downarrow & \uparrow & \downarrow & \downarrow & \downarrow \\ \uparrow & \uparrow & \downarrow & \uparrow & \downarrow \\ \uparrow & \downarrow & \uparrow & \uparrow & \downarrow \end{array} \sim e^{-\frac{E}{k_B T}}$$

- **Algorithm:**

1.) Propose new state: Random spin flips!

$$\begin{array}{ccccc} \uparrow & \downarrow & \downarrow & \uparrow & \uparrow \\ \downarrow & \uparrow & \downarrow & \uparrow & \uparrow \\ \uparrow & \uparrow & \uparrow & \downarrow & \downarrow \end{array} \Rightarrow \begin{array}{ccccc} \uparrow & \downarrow & \downarrow & \uparrow & \uparrow \\ \downarrow & \uparrow & \uparrow & \uparrow & \uparrow \\ \uparrow & \uparrow & \uparrow & \downarrow & \downarrow \end{array}$$


2.) Accept the new configuration:  $p_a = e^{-\frac{\Delta E}{k_B T}}$

## Single core CPU: Data structure

- ▶ Multi-spin coding:  $1 \text{ spin} \hat{=} 1 \text{ bit}$


## Single core CPU: Data structure

- ▶ Multi-spin coding: 1 spin  $\hat{=}$  1 bit
- ▶ Group spins in groups of size 32 (int)




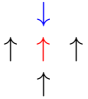
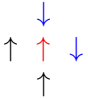
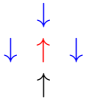
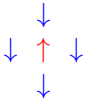
## Single core CPU: Data structure

- ▶ Multi-spin coding: 1 spin  $\hat{=}$  1 bit
- ▶ Group spins in groups of size 32 (int)

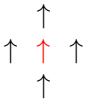
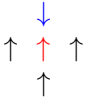
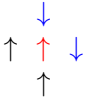
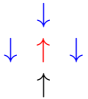
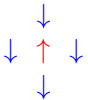


- ▶ 100'000 × 100'000 lattice:  $\approx$  1.2 GB

# Single core CPU: Algorithm I

Type:	# opposed	$\Delta E$ caused by flip of $\uparrow$
	0	$+8J$
	1	$+4J$
	2	0
	3	$-4J$
	4	$-8J$

# Single core CPU: Algorithm I

Type:	# opposed	$\Delta E$ caused by flip of $\uparrow$
	0	$+8J$
	1	$+4J$
	2	0
	3	$-4J$
	4	$-8J$

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

# Single core CPU: Algorithm I

Type:	# opposed	$\Delta E$ caused by flip of $\uparrow$
$\begin{array}{c} \uparrow \\ \uparrow \uparrow \uparrow \\ \uparrow \end{array}$	0	$+8J$
$\begin{array}{c} \downarrow \\ \uparrow \uparrow \uparrow \\ \uparrow \end{array}$	1	$+4J$
$\begin{array}{c} \downarrow \\ \uparrow \uparrow \downarrow \\ \uparrow \end{array}$	2	0
$\begin{array}{c} \downarrow \\ \downarrow \uparrow \downarrow \\ \uparrow \end{array}$	3	$-4J$
$\begin{array}{c} \downarrow \\ \downarrow \uparrow \downarrow \\ \downarrow \end{array}$	4	$-8J$

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$

$$p_a = e^{-\frac{\Delta E}{k_B T}}$$



# Single core CPU: Algorithm I


Type:	# opposed	$\Delta E$ caused by flip of $\uparrow$
$\begin{array}{c} \uparrow \\ \uparrow \uparrow \uparrow \\ \uparrow \end{array}$	0	$+8J$
$\begin{array}{c} \downarrow \\ \uparrow \uparrow \uparrow \\ \uparrow \end{array}$	1	$+4J$
$\begin{array}{c} \downarrow \\ \uparrow \uparrow \downarrow \\ \uparrow \end{array}$	2	0
$\begin{array}{c} \downarrow \\ \downarrow \uparrow \downarrow \\ \uparrow \end{array}$	3	$-4J$
$\begin{array}{c} \downarrow \\ \downarrow \uparrow \downarrow \\ \downarrow \end{array}$	4	$-8J$

$$H = -J \sum_{\langle i,j \rangle} S_i S_j$$


$$p_a = e^{-\frac{\Delta E}{k_B T}}$$

Count  
spins!      #opposed

## Single core CPU: Algorithm II




## Single core CPU: Algorithm II



$$\text{red box} \neq \text{green box} \Leftrightarrow \text{red box XOR green box}$$

## Single core CPU: Algorithm II



$$\text{red box} \neq \text{green box} \Leftrightarrow \text{red box XOR green box}$$


$$i_1 = \text{red box XOR green box}$$

$$i_2 = \text{red box XOR yellow box}$$

$$i_3 = \text{red box XOR blue box}$$

$$i_4 = \text{red box XOR black box}$$

## Single core CPU: Algorithm II



$$\text{red box} \neq \text{green box} \Leftrightarrow \text{red box XOR green box}$$

$$i_1 = \text{red box XOR green box}$$

$$i_2 = \text{red box XOR yellow box}$$

$$i_3 = \text{red box XOR blue box}$$

$$i_4 = \text{red box XOR black box}$$

Combine with acceptance probability:

$$i_1 + i_2 + i_3 + i_4 + 2\text{exp}_8 + \text{exp}_4 \geq 2$$

# Single GPU implementation I

- ▶ Port CPU implementation

# Single GPU implementation I

- ▶ Port CPU implementation  
☹ Bad performance!

# Single GPU implementation I

- ▶ Port CPU implementation  
☹ Bad performance!
- ▶ GPU in a nutshell:



# Single GPU implementation I

- ▶ Port CPU implementation  
☹ Bad performance!
- ▶ GPU in a nutshell:
  - ▶ 100s of cores, blocks of 512

# Single GPU implementation I

- ▶ Port CPU implementation
  - ☹ Bad performance!
- ▶ GPU in a nutshell:
  - ▶ 100s of cores, blocks of 512
  - ▶ Memory:

# Single GPU implementation I

- ▶ Port CPU implementation
  - ☹ Bad performance!
- ▶ GPU in a nutshell:
  - ▶ 100s of cores, blocks of 512
  - ▶ Memory:  
Global memory:



- ▶ big ( $\approx 4$  GB)
- ▶ slow


# Single GPU implementation I

- ▶ Port CPU implementation
  - ☹ Bad performance!
- ▶ GPU in a nutshell:
  - ▶ 100s of cores, blocks of 512
  - ▶ Memory:  
Global memory:



- ▶ big ( $\approx 4$  GB)
- ▶ slow

Shared memory:



- ▶ small ( $\approx 16$  kB per block)
- ▶ fast


# Single GPU implementation I

- ▶ Port CPU implementation
  - ☹ Bad performance!
- ▶ GPU in a nutshell:
  - ▶ 100s of cores, blocks of 512
  - ▶ Memory:
    - Global memory:



- ▶ big ( $\approx 4$  GB)
- ▶ slow

Shared memory:



- ▶ small ( $\approx 16$  kB per block)
- ▶ fast


Reduce # accesses to global memory!

## Single GPU implementation II

- ▶ Metaspins ( $4 \times 4$ )
  - 1 metaspin  $\hat{=}$  2 bytes = 1 unsigned short int
  - only 1 read per metaspin


# Single GPU implementation II

- ▶ Metaspins ( $4 \times 4$ )
  - 1 metaspin  $\hat{=}$  2 bytes = 1 unsigned short int
  - only 1 read per metaspin
- ▶ Read:




## Single GPU implementation II

Global memory:



5reads  
→


Shared memory:



⇒ 5 reads to flip entire metaspin!



## Results: CPU vs. GPU



# Multi-GPU approach

- ▶ Single GPU:  
fast

# Multi-GPU approach

- ▶ Single GPU:

fast

system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )

# Multi-GPU approach

- ▶ Single GPU:  
fast  
system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )
- ▶ Idea: Distributed lattice!

# Multi-GPU approach

- ▶ Single GPU:  
fast  
system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )
- ▶ Idea: Distributed lattice!
- ▶ Algorithm:
  1. Copy neighbour borders to GPU

# Multi-GPU approach

- ▶ Single GPU:  
fast  
system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )
- ▶ Idea: Distributed lattice!
- ▶ Algorithm:
  1. Copy neighbour borders to GPU
  2. Update own region on GPU

# Multi-GPU approach

- ▶ Single GPU:  
fast  
system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )
- ▶ Idea: Distributed lattice!
- ▶ Algorithm:
  1. Copy neighbour borders to GPU
  2. Update own region on GPU
  3. Copy boundary spins to CPU

# Multi-GPU approach

- ▶ Single GPU:  
fast  
system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )
- ▶ Idea: Distributed lattice!
- ▶ Algorithm:
  1. Copy neighbour borders to GPU
  2. Update own region on GPU
  3. Copy boundary spins to CPU
  4. Exchange boundary spins with other nodes



# Multi-GPU approach

- ▶ Single GPU:

fast


system size  $\leq 4$  GB ( $\hat{=}$   $100'000 \times 100'000$ )

- ▶ Idea: Distributed lattice!


- ▶ Algorithm:

1. Copy neighbour borders to GPU
2. Update own region on GPU
3. Copy boundary spins to CPU
4. Exchange boundary spins with other nodes
5. Repeat or finish

## Multi-GPU: Results



## Multi-GPU: Results



64 GPUs, 800'000 × 800'000: 3s

# Conclusion

# Conclusion

- ▶ No uncertainties in benchmark plots

# Conclusion

- ▶ No uncertainties in benchmark plots
- ▶ Why this system?

# Conclusion

- ▶ No uncertainties in benchmark plots
- ▶ Why this system?
- ▶ Approach generalisable?

# Conclusion

- ▶ No uncertainties in benchmark plots
- ▶ Why this system?
- ▶ Approach generalisable?
- ▶ Not examined big systems (no weak scaling plot)



# Conclusion

- ▶ No uncertainties in benchmark plots
- ▶ Why this system?
- ▶ Approach generalisable?
- ▶ Not examined big systems (no weak scaling plot)

Why this paper?!

**Questions?**